


Heap Sort is an algo that relies on a specific data structure called a *heap*.

Heaps can be top-heavy, called a max heap, or bottom-heavy, called a min heap.




Heaps are binary trees: each node has zero, one, or two children. In a max heap a parent node is at least as large as its children. No distinction is made between left and right branches. Min heaps have smaller values at the top instead.


The first step in using heapsort is to build a heap out of your data. E.g. given the list:

[10, 15, 2, 4, 17, 12]

This is one possible
resulting heap.




When we have a heap like this, we say that the *order* of the heap is like this: first node, then second row (left to right), then third row (left to right), and so on.




Heap order:
17, 15, 12, 4, 2, 10

The algorithm now is pretty straight forward.

1. Swap the top of the heap with the last element.
2. Remove the last element, add it to sorted list.
3. Repair the heap: large numbers above small ones.
4. Repeat until heap has nothing in it.




At start we swap the top node (17) with the 'last' node: bottom row, farthest to the right.



Return list:

(empty)


Now 'remove' the 17 from the heap. You can do this in the obvious way, or you can keep it there but just remember it is no longer in the heap.



Return list:

(empty)

Notice that the heap
order is messed up.
The top node is smaller
than its children!
Fix that.




*After removing
end node.*

Return list:

17


We fix the heap by
swapping the 'wrong'
node with one of
the children.



Return list:

17


Only needed one swap.
In general we may
need to do this a
bunch of times.
Up to the tree
height.



Return list:

17


Back to the start of the algorithm. Swap top node with the 'last' node.



Return list:

17

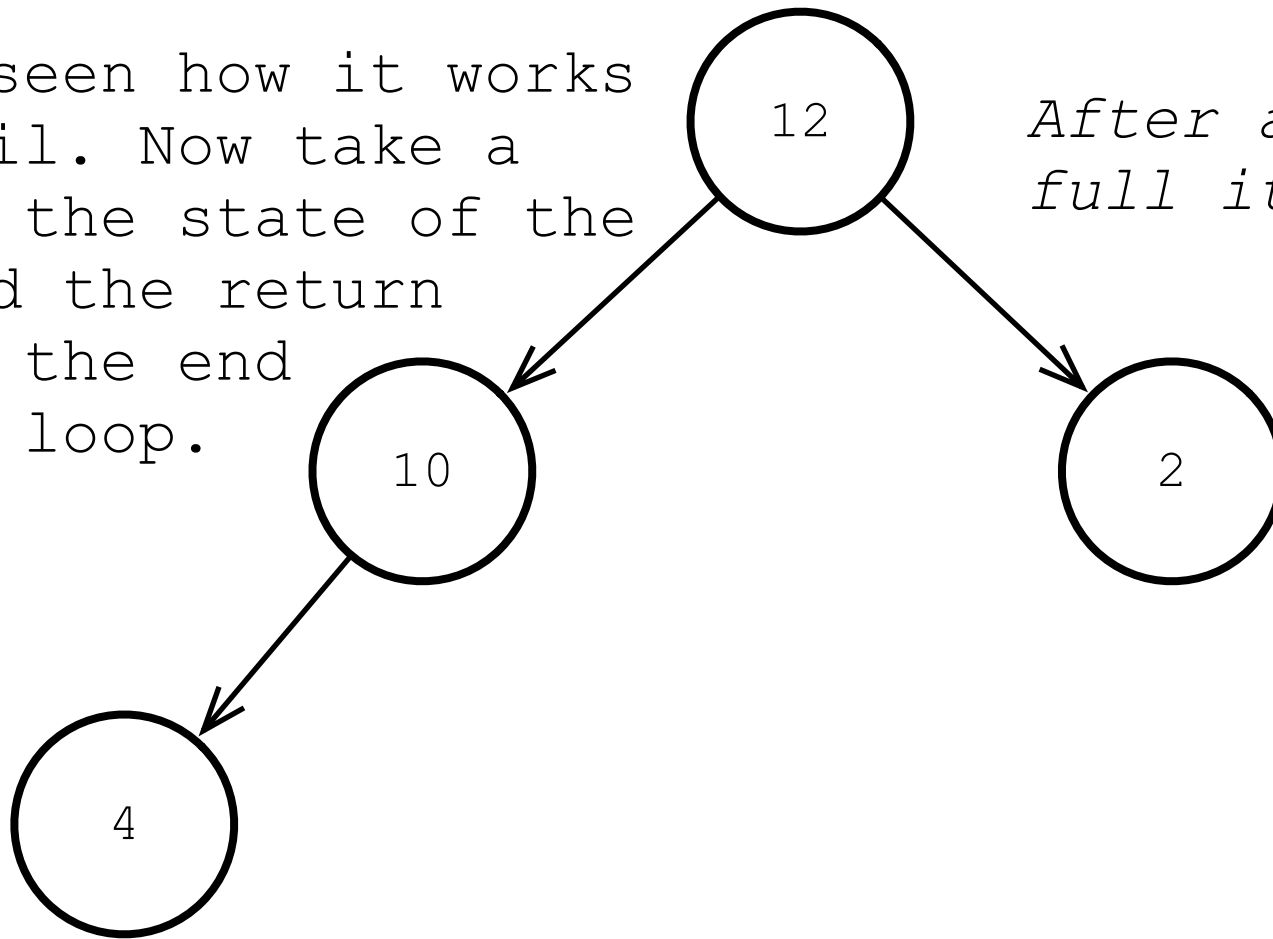
Swapped again. 15 has
to be removed and heap
has to be repaired.



Return list:

17


You've seen how it works in detail. Now take a look at the state of the heap and the return list at the end of each loop.



After another full iteration.


Return list:

17, 15




Return list:

17, 15, 12



Return list:

17, 15, 12, 10



Return list:

17, 15, 12, 10, 4

Finally at the end. Notice that the return list is actually in the reverse order from what we expect. We could fix this in a couple ways. We could use a *min heap* instead of a *max heap*.

Return list:

17, 15, 12, 10, 4, 2

Or, we could use a trick with the tree structure. If we swap with the last node and then *keep it around*, but just make it off limits to the rest of the algorithm, the resulting heap structure is in proper sort order.

For added brain melt, you can implement this without even using a binary tree. You can use an array instead. I'll just leave it at that :)

